

# Greedy Algorithms

---

Atli FF

September 9, 2024

School of Computer Science

Reykjavík University

# Greedy algorithms

- An algorithm that always makes *locally* optimal moves is called greedy
- For some kinds of problems this will give a *globally* optimal solution as well
- Seeing when this is the case can be very tricky, and if used in the wrong context the solution will get a WA verdict

## Submitting greedy solutions

- The tricky thing with these solutions are that it's often hard to know if you've made a mistake and thus get WA or if there's some hole in the greedy algorithm
- It's often easy to think of all kinds of greedy solutions, but they are very often wrong
- Generally one would like to consider complete search or dynamic programming (will see this later) first, but some problems do require greedy solutions

# Coin change

- A classical example is making change. Say you want to sum up  $n$  and have only denominations of 1, 5 and 10, what's the least amount of coins you can give back?
- The greedy solution would be to just always give the biggest coin you can that's not too much. So for say 24 we'd do 10, 10, 1, 1, 1, 1.
- Is this always optimal?

# Coin change

- Well, it turns out to depend on the denominations. Say we have denominations of 1, 8 and 20.
- For  $n = 24$  we then give back 20, 1, 1, 1, 1 instead of the optimal 8, 8, 8.
- We will come back to this problem when we solve the general case using dynamic programming.

# Lilypad jump

- Consider a frog jumping on a sequence of lily pads, there is one at  $x = 0$  and one at  $x = n$ , with some amount of lily pads in between
- The frog can jump at most distance  $r$
- When at a given lily pad, what's the best move?

# Lily pad jump

- Clearly just jump as far right as possible!
- But be careful, this is very contingent on the frog being able to jump any distance in  $[0, r]$
- If it could jump any distance in  $[r/2, r]$ , it would not be true for example

# Taxi assignment

- Let's consider another problem. You are managing a taxi company and today  $n$  drivers showed up and you have  $m$  cars.
- But not all drivers and cars are created equal. Car  $i$  has  $h_i$  horsepower and driver  $j$  can only handle at most  $g_j$  horsepower.
- What's the greatest number of drivers you can pair to cars such that they can handle their car?



## The greedy step

- The greedy idea here is to simply pair each car to the worst driver that can still handle that car.
- Thus we start by sorting the drives and cars and then simply linearly walk through each and pair them together.
- It might not be obvious, but this actually gives the best answer.

# Implementation

```
int main() {  
    int n, m; cin >> n >> m;  
    vi a(n), b(m);  
    for(int i = 0; i < n; ++i) cin >> a[i];  
    for(int i = 0; i < m; ++i) cin >> b[i];  
    sort(a.begin(), a.end());  
    sort(b.begin(), b.end());  
    int ans = 0;  
    for(int i = 0, j = 0; i < m; ++i) {  
        while(j < n && a[j] < b[i]) j++;  
        if(j < n) ans++, j++;  
    }  
    cout << ans << '\n';  
}
```

# Sorting

- Greedy algorithms very often involve sorting
- More generally they often involve always picking the “extremal” option out of the local options, in some sense
- Biggest, shortest, cheapest, first, etc.

# Job scheduling

- Say we have a list of jobs, each starting at some time  $s_j$  and finishing at some time  $f_j$
- What's the largest amount of jobs we can complete if they can't overlap?

# Solution

- The solution is shockingly simple, but not obviously correct
- Order the jobs by completion time  $f_j$  and then walk through them
- If you can complete a job in addition to the ones you've already picked, pick it
- The jobs you've picked by the end are the solution

## Proof of correctness

- Why is this correct though? Let's prove it.
- Suppose the algorithm is not optimal. Say we pick jobs of indices  $i_1, i_2, \dots, i_k$  but a better solution picks  $j_1, j_2, \dots, j_l$ .
- Say the solutions agree on the first  $r$  jobs (possibly 0).
- Now neither  $i_{r+1}$  nor  $j_{r+1}$  clash with the jobs  $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ . But because we ordered things by end time, we must have that job  $i_{r+1}$  ends no later than  $j_{r+1}$ . But then we could just as well have picked  $i_{r+1}$ . But this holds for any  $r$ , so by induction we have that  $i_1, \dots, i_k$  is no worse than  $j_1, \dots, j_l$ , which gives a contradiction.
- Thus the algorithm is optimal.

# Many more

- There are many many more and we will see plenty in the course
- Many famous algorithms are famous because they perform non-trivial greedy steps
- Dijkstra's algorithm, Huffman coding, Kruskal's algorithm, Horn satisfiability and many more

# Proving correctness

- How do we prove the greedy algorithms are correct?
- When in programming contests this is usually overkill, but at a workplace it is generally not
- There are two main common arguments that tackle this, but often novel methods are needed
- These are usually called "Greedy stays ahead" and "Exchange arguments"



# Proving correctness

- "Greedy stays ahead" aims to prove that during each greedy step it consistently stays ahead of all other possible choices
- "Exchange arguments" aims to show that you can turn any solution into the greedy solution with a sequence of "exchanges" without making them any worse, so the greedy solution is as good as it gets
- Sometimes a part of this is showing that the greedy solution outputs a valid solution at all, as that is not obvious in every case

# Proving correctness

- Consider the lilypads again.
- We'll now prove it's optimal, which is more work than it may seem at first
- First we prove that the greedy solution reaches the final lily pad if there is a path there
- Note again this is not true for many minor variants of the problem!

## Proving correctness

By contradiction; suppose it did not. Let the positions of the lilypads be  $x_1 < x_2 < \dots < x_m$ . Since our algorithm didn't find a path, it must have stopped at some lilypad  $x_k$  and not been able to jump to a future lilypad. In particular, this means it could not jump to lilypad  $k + 1$ , so  $x_k + r < x_{k+1}$ . Since there is a path from lilypad 1 to the lilypad  $m$ , there must be some jump in that path that starts before lilypad  $k + 1$  and ends at or after lilypad  $k + 1$ . This jump can't be made from lilypad  $k$ , so it must have been made from lilypad  $s$  for some  $s < k$ . But then we have  $x_s + r < x_k + r < x_{k+1}$ , so this jump is illegal. We have reached a contradiction, so our assumption was wrong and our algorithm always finds a path.

# Proving correctness

- Let's now show it actually stays ahead of any optimal solution
- Let  $J$  be the set of jumps from our greedy algorithm and  $J^*$  be an optimal set of jumps
- Then  $|J| \geq |J^*|$  since it's optimal
- Let  $p(i, J)$  be the position after taking the first  $i$  jumps in  $J$
- Let's prove that for all  $i$  we have  $p(i, J) \geq p(i, J^*)$

## Proving correctness

We proceed by induction. As a base case, if  $i = 0$ , then  $p(0, J) = 0 \geq 0 = p(0, J^*)$  since the frog hasn't moved. For the inductive step, assume that the claim holds for some  $0 \leq i < |J^*|$ . We will prove the claim holds for  $i + 1$  by considering two cases:

- Case 1:  $p(i, J) \geq p(i + 1, J^*)$ . Since each jump moves forward, we have  $p(i + 1, J) \geq p(i, J)$ , so we have  $p(i + 1, J) \geq p(i + 1, J^*)$ .
- Case 2:  $p(i, J) < p(i + 1, J^*)$ . Each jump is of size at most  $r$ , so  $p(i + 1, J^*) \leq p(i, J^*) + r$ . By the inductive hypothesis, we know  $p(i, J) \geq p(i, J^*)$ , so  $p(i + 1, J^*) \leq p(i, J) + r$ . Therefore, the greedy algorithm can jump to position at least  $p(i + 1, J^*)$ . Therefore,  $p(i + 1, J) \geq p(i + 1, J^*)$ .

So  $p(i + 1, J) \geq p(i + 1, J^*)$ , completing the induction.

# Proving correctness

- And now we are almost done!
- Finally we just have to prove that  $|J| = |J^*|$ , which would mean the greedy is always optimal

## Proving correctness

Since  $J^*$  is an optimal solution, we know that  $|J^*| \leq |J|$ . We will prove  $|J^*| \geq |J|$ . Suppose for contradiction that  $|J^*| < |J|$ . Let  $k = |J^*|$ . From before, we have  $p(k, J^*) \leq p(k, J)$ . Because the frog arrives at position  $n$  after  $k$  jumps along series  $J^*$ , we know  $n \leq p(k, J)$ . Because the greedy algorithm never jumps past position  $n$ , we know  $p(k, J) \leq n$ , so  $n = p(k, J)$ . Since  $|J^*| < |J|$ , the greedy algorithm must have taken another jump after its  $k$ -th jump, contradicting that the algorithm stops after reaching position  $n$ . We have reached a contradiction, so our assumption was wrong and  $|J^*| = |J|$ , so the greedy algorithm produces an optimal solution.

# Proving correctness

- That was a lot of effort!
- Now, which kind of greedy proof was that?



# Proving correctness

- That was a lot of effort!
- Now, which kind of greedy proof was that?
- It was a "Greedy stays ahead" proof, so let us next see an exchange proof

# Proving correctness

- To consider exchange arguments we have to do things slightly out of order and nab an algorithm from the future week of graph theory
- Consider a set of houses, we want to lay fibre cable between them such that each house is connected to every other house through some set of cables
- Doesn't have to be a direct connection, just some path exists between them
- For each pair of houses we are given the cost of laying a cable between them
- What's the cheapest cable-laying procedure?

# Proving correctness

- Our greedy algorithm will be as follows
- Start with just a single house and consider it "active"
- Choose the cheapest cable that goes between an active house and one that is not active
- Make the newly connected house active, and keep going
- Now we prove this is optimal!

## Proving correctness

Let  $T$  be the set of cables chosen by our algorithm and  $T^*$  be some optimal set. Let  $c(T)$  denote the total cost of a set of cables. We will prove  $c(T) = c(T^*)$ . If  $T = T^*$  the result is obvious, so we can assume  $T \neq T^*$ . Then let  $(u, v)$  be a cable in  $T \setminus T^*$ . Let  $S$  be the set of active houses when  $(u, v)$  was added to  $T$  and  $H$  be the set of all houses. Then  $(u, v)$  is the cheapest edge between  $S$  and  $H \setminus S$ . Since  $T^*$  connects all houses it must contain some path from  $u$  to  $v$ . This path begins in  $S$  and ends in  $H \setminus S$  so there must be some cable  $(x, y)$  such that  $x \in S$  and  $y \in H \setminus S$ . Since  $(u, v)$  is the cheapest such edge we must have  $c(\{(u, v)\}) \leq c(\{(x, y)\})$ . Let  $T' = T^* \cup \{(u, v)\} \setminus \{(x, y)\}$ .

## Proving correctness

Since every house in  $S$  can reach every other house in  $S$  without using  $(u, v)$ ,  $T'$  is valid for those houses, and the same goes for  $H \setminus S$ . But then  $T'$  allows any house in  $S$  to reach  $u$ , then go to  $v$ , then to any house in  $H \setminus S$ . So  $T'$  is a valid set of connections. But note that  $c(T') = c(T^*) - c(\{(x, y)\}) + c(\{(u, v)\}) \leq c(T^*)$ . Since  $T^*$  is optimal this means  $c(T') \geq c(T^*)$ , so  $c(T') = c(T^*)$ . Note that  $|T \setminus T'| = |T \setminus T^*| - 1$ , so if we repeat this same argument once for each edge in  $T \setminus T^*$  we will have converted  $T^*$  into  $T$  without changing  $T$ , thus  $c(T) = c(T^*)$ .

# Proving correctness

- Very mathy!
- But that is an example of an "Exchange argument" proof of correctness