

# Greedy Algorithms

---

Atli FF

September 8, 2024

School of Computer Science

Reykjavík University

# Greedy algorithms

- An algorithm that always makes *locally* optimal moves is called greedy
- For some kinds of problems this will give a *globally* optimal solution as well
- Seeing when this is the case can be very tricky, and if used in the wrong context the solution will get a WA verdict

# Submitting greedy solutions

- The tricky thing with these solutions are that it's often hard to know if you've made a mistake and thus get WA or if there's some hole in the greedy algorithm
- It's often easy to think of all kinds of greedy solutions, but they are very often wrong
- Generally one would like to consider complete search or dynamic programming (will see this later) first, but some problems do require greedy solutions

# Coin change

- A classical example is making change. Say you want to sum up  $n$  and have only denominations of 1, 5 and 10, what's the least amount of coins you can give back?
- The greedy solution would be to just always give the biggest coin you can that's not too much. So for say 24 we'd do 10, 10, 1, 1, 1, 1.
- Is this always optimal?

# Coin change

- Well, it turns out to depend on the denominations. Say we have denominations of 1, 8 and 20.
- For  $n = 24$  we then give back 20, 1, 1, 1, 1 instead of the optimal 8, 8, 8.
- We will come back to this problem when we solve the general case using dynamic programming.

# Taxi assignment

- Let's consider another problem. You are managing a taxi company and today  $n$  drivers showed up and you have  $m$  cars.
- But not all drivers and cars are created equal. Car  $i$  has  $h_i$  horsepower and driver  $j$  can only handle at most  $g_j$  horsepower.
- What's the greatest number of drivers you can pair to cars such that they can handle their car?

# The greedy step

- The greedy idea here is to simply pair each car to the worst driver that can still handle that car.
- Thus we start by sorting the drives and cars and then simply linearly walk through each and pair them together.
- It might not be obvious, but this actually gives the best answer.

# Implementation

```
int main() {  
    int n, m; cin >> n >> m;  
    vi a(n), b(m);  
    for(int i = 0; i < n; ++i) cin >> a[i];  
    for(int i = 0; i < m; ++i) cin >> b[i];  
    sort(a.begin(), a.end());  
    sort(b.begin(), b.end());  
    int ans = 0;  
    for(int i = 0, j = 0; i < m; ++i) {  
        while(j < n && a[j] < b[i]) j++;  
        if(j < n) ans++, j++;  
    }  
    cout << ans << '\n';  
}
```



# Sorting

- Greedy algorithms very often involve sorting
- More generally they often involve always picking the “extremal” option out of the local options, in some sense
- Biggest, shortest, cheapest, first, etc.

# Job scheduling

- Say we have a list of jobs, each starting at some time  $s_j$  and finishing at some time  $f_j$
- What's the largest amount of jobs we can complete if they can't overlap?

# Solution

- The solution is shockingly simple, but not obviously correct
- Order the jobs by completion time  $f_j$  and then walk through them
- If you can complete a job in addition to the ones you've already picked, pick it
- The jobs you've picked by the end are the solution

## Proof of correctness

- Why is this correct though? Let's prove it.
- Suppose the algorithm is not optimal. Say we pick jobs of indices  $i_1, i_2, \dots, i_k$  but a better solution picks  $j_1, j_2, \dots, j_l$ .
- Say the solutions agree on the first  $r$  jobs (possibly 0).
- Now neither  $i_{r+1}$  nor  $j_{r+1}$  clash with the jobs  $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ . But because we ordered things by end time, we must have that job  $i_{r+1}$  ends no later than  $j_{r+1}$ . But then we could just as well have picked  $i_{r+1}$ . But this holds for any  $r$ , so by induction we have that  $i_1, \dots, i_k$  is no worse than  $j_1, \dots, j_l$ , which gives a contradiction.
- Thus the algorithm is optimal.