

# Graphs Part 2

---

Atli Fannar Franklín

October 10, 2025

School of Computer Science

Reykjavík University

# Today we're going to cover

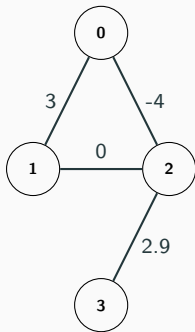
- Minimum spanning tree
- Shortest paths
  - Single source
  - Negative weights
  - All pairs

# Weighted graphs

- Now the edges in our graphs may have weights, which could represent
  - the distance of the road represented by the edge
  - the cost of going over the edge
  - some capacity of the edge
- We can use a modified adjacency list to represent weighted graphs

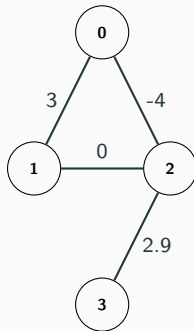
# Weighted graphs

```
struct edge {  
    // Depending on context,  
    // having both endpoints  
    // can be overkill  
    int u, v;  
    int weight;  
    // Sometimes weight has to  
    // be a double, or something else  
  
    edge(int _u, int _v, int _w) {  
        u = _u;  
        v = _v;  
        weight = _w;  
    }  
};
```



# Weighted graphs

```
vector<edge> adj[4];  
  
adj[0].push_back(edge(0, 1, 3));  
adj[0].push_back(edge(0, 2, -4));  
  
adj[1].push_back(edge(1, 0, 3));  
adj[1].push_back(edge(1, 2, 0));  
  
adj[2].push_back(edge(2, 0, -4));  
adj[2].push_back(edge(2, 1, 0));  
adj[2].push_back(edge(2, 3, 2.9));  
  
adj[3].push_back(edge(3, 2, 2.9));
```



# Minimum spanning trees

---

# Minimum spanning tree

- We have an undirected weighted graph
- The vertices along with a subset of the edges in the graph is called a spanning tree if
  - it forms a tree (i.e. does not contain a cycle) and
  - the tree spans all vertices (all vertices can reach all other vertices)
- The weight of a spanning tree is the sum of the weights of the edges in the subset
- We want to find a minimum spanning tree

# Minimum spanning tree

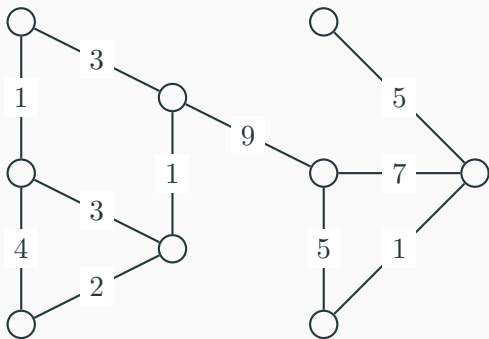
- This can model many things, the cheapest way to connect a set of objects is one of them
- For some graph problems we can remove everything but a spanning tree to simplify calculations
- Note however that it does not have to be unique, but the minimum weight is unique



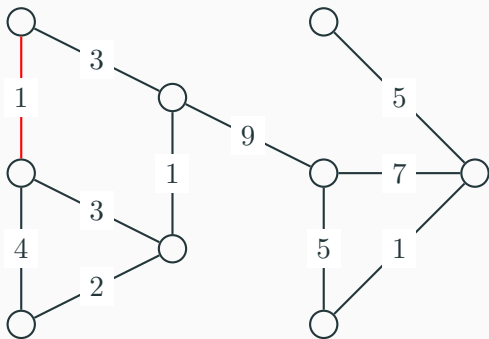
# Minimum spanning tree

- Several greedy algorithms work
- Go through the edges in the graph in increasing order of weight
- Greedily pick an edge if it doesn't form a cycle (Union-Find can be used to keep track of when we would get a cycle)
- When we've gone through all edges, we have a minimum spanning tree
- This is Kruskal's algorithm
- Time complexity is  $O(E \log E)$

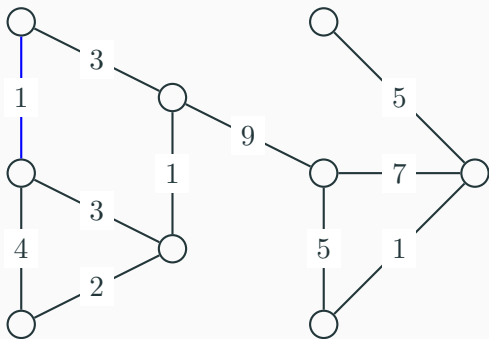
## Example



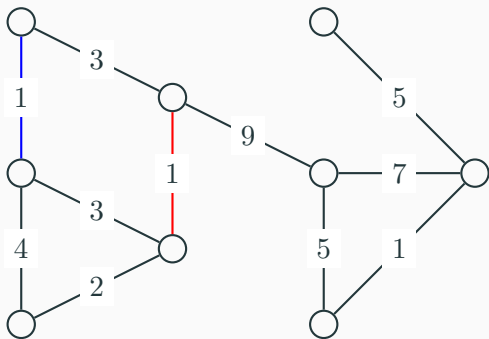
## Example



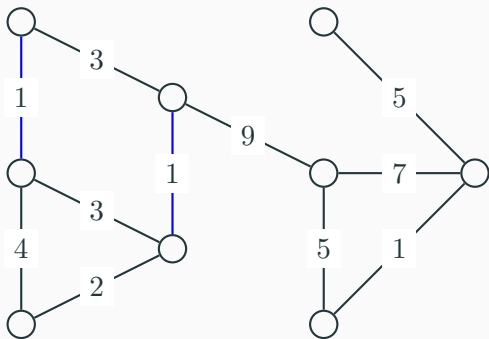
## Example



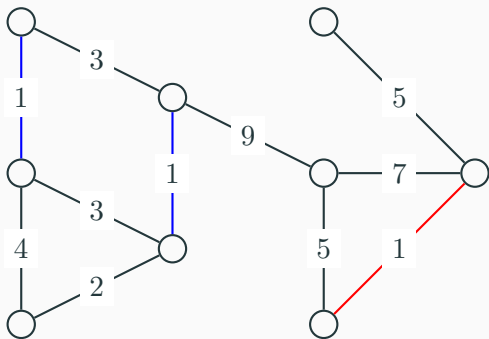
## Example



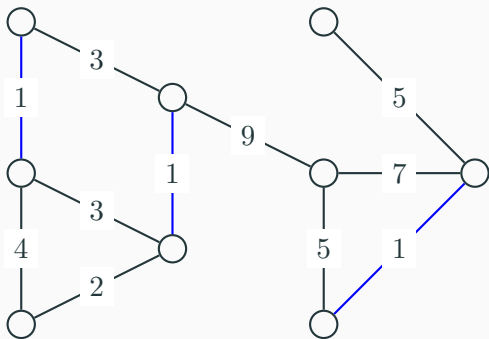
## Example



## Example

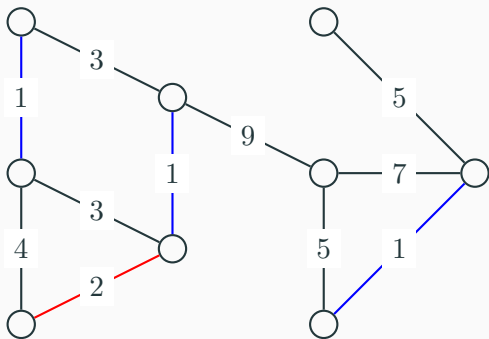


## Example

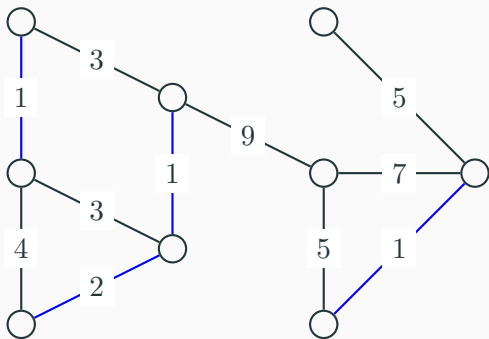




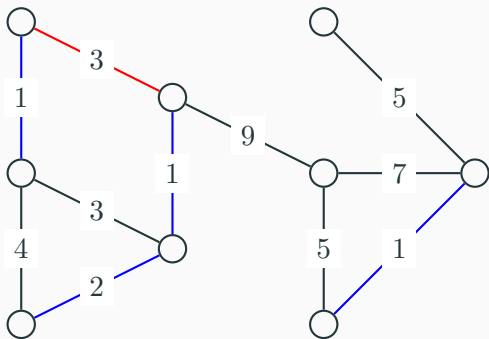
## Example



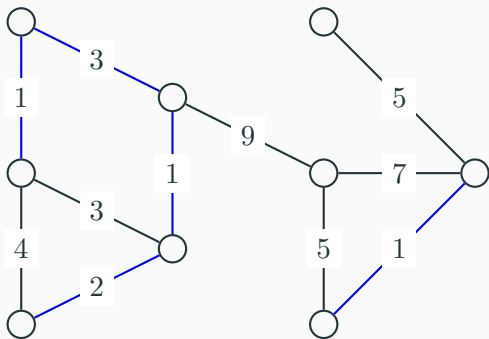
## Example



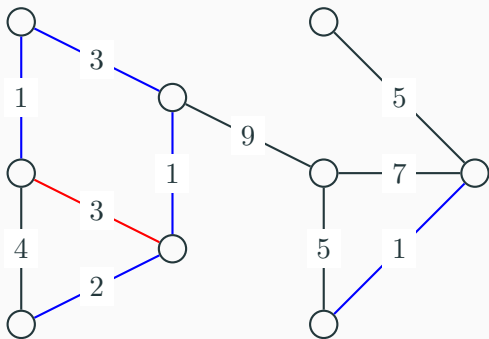
## Example



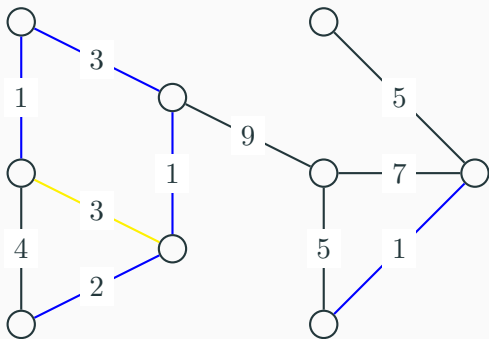
## Example



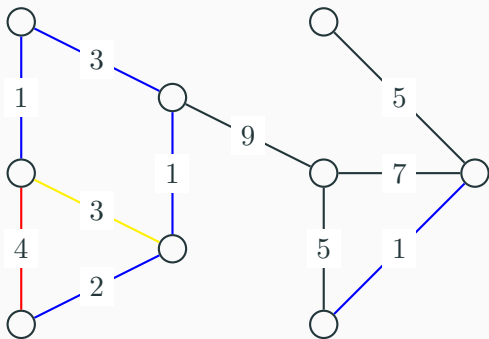
## Example



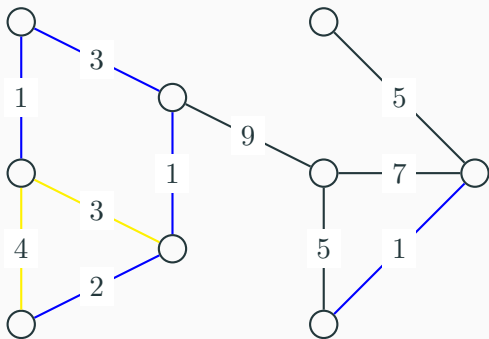
## Example



## Example

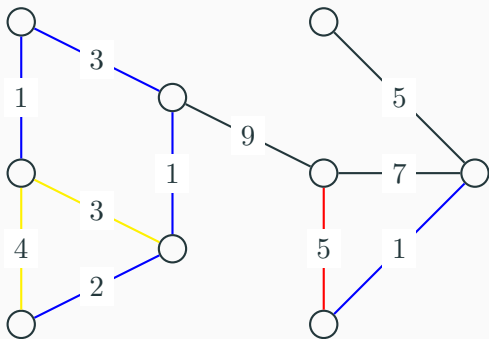


## Example

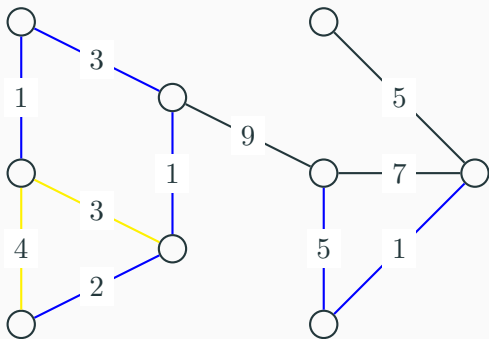




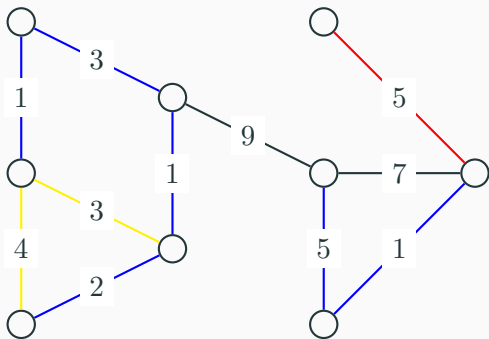
## Example



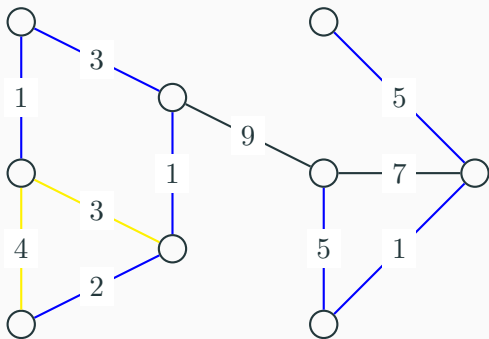
## Example



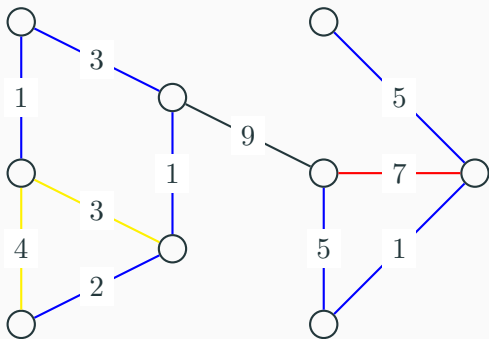
## Example



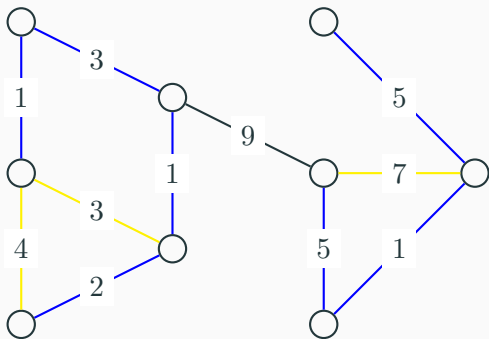
## Example



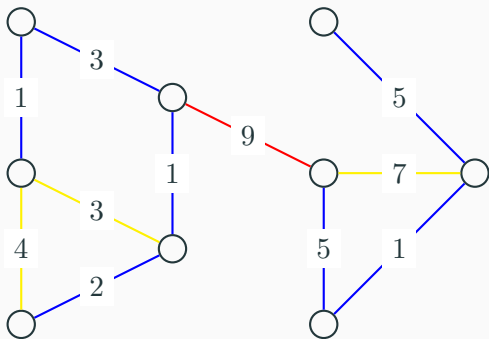
## Example



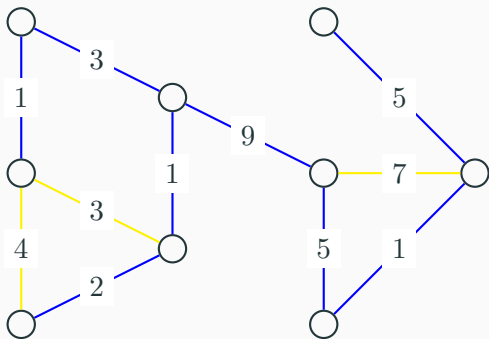
## Example



## Example




## Example





## Example



# Minimum spanning tree

```
bool edge_cmp(const edge &a, const edge &b) {
    return a.weight < b.weight;
}

vector<edge> mst(int n, vector<edge> edges) {
    union_find uf(n);
    sort(edges.begin(), edges.end(), edge_cmp);

    vector<edge> res;
    for (int i = 0; i < edges.size(); i++) {
        int u = edges[i].u,
            v = edges[i].v;

        if (uf.find(u) != uf.find(v)) {
            uf.unite(u, v);
            res.push_back(edges[i]);
        }
    }

    return res;
}
```

# Minimum spanning tree

- Another one that works is Prim's algorithm
- We start with vertex 0, then grow the tree one vertex at a time
- Each time we take the cheapest edge that connects our tree to something not in the tree
- This can also be done in  $\mathcal{O}(E \log(E))$
- For our purposes Kruskal will suffice

# Single source shortest path

---

# Shortest paths

- We have a weighted graph (undirected or directed)
- Given two vertices  $u, v$ , what is the shortest path from  $u$  to  $v$ ?
- If all weights are the same (or 0/1), this can be solved with breadth-first search
- Of course, this is usually not the case...

# Shortest paths

- There are many known algorithms to find shortest paths
- Like breadth-first search, these algorithms usually find the shortest paths from a given start vertex to all other vertices
- We will look at Dijkstra's algorithm, the Bellman-Ford algorithm and the Floyd-Warshall algorithm

# Shortest paths

- We will start with the simplest case, where no weights are negative and we want the distance from a particular vertex to the other ones
- Why no negative weights? Well if there is a cycle with negative weight on the way to our destination there is no shortest path! You can always go a "shorter" way by cycling around the negative cycle an extra time
- In principle Dijkstra is not so dissimilar from BFS

# Shortest paths


- In BFS the next item in our queue was always the closest one to the source vertex we had not processed yet
- In Dijkstra's algorithm we have to do some extra work to ensure this is the case
- We use a min-heap, so that the next vertex we pick is always the closest unprocessed one



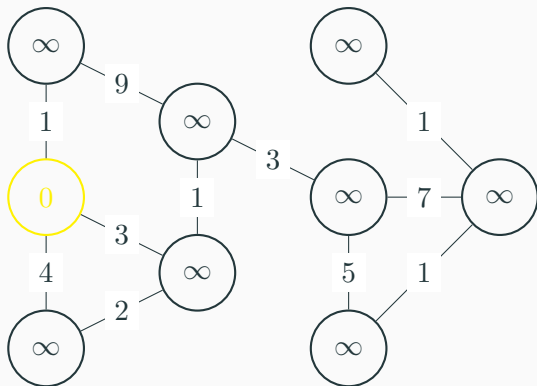
# Shortest paths

- We initialize the source vertex to distance 0, others to  $\infty$
- We then at any point take the closest vertex  $v$  that we haven't processed
- We then check all edges adjacent to  $v$  and update the distance to the neighbours
- We then mark  $v$  as done and don't check it again
- Note that if all weights are one this will just do the same things as BFS

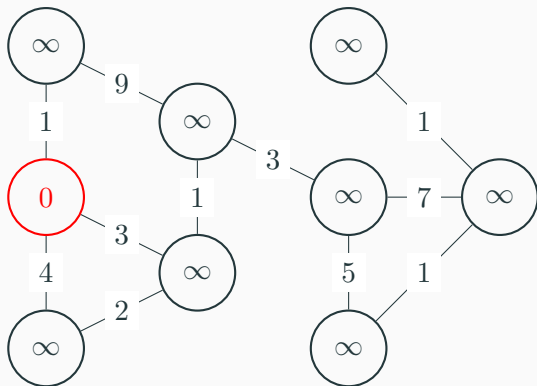
## Example



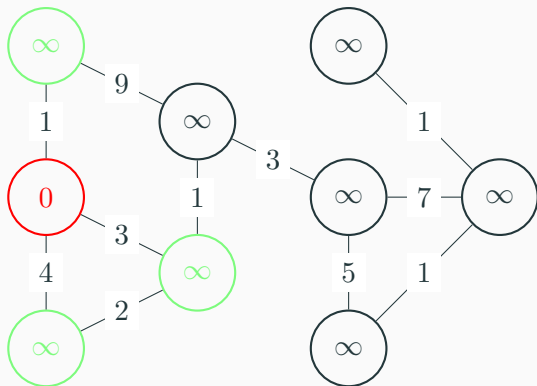
## Example



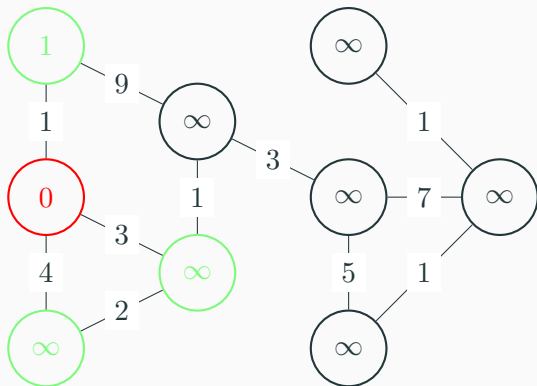
## Example



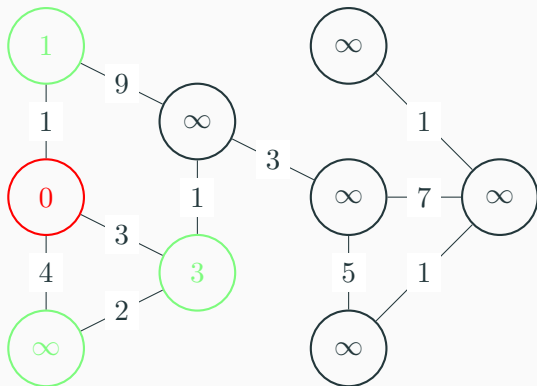
# Example



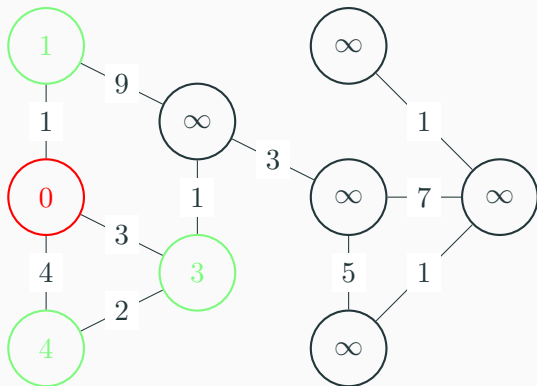
## Example



## Example

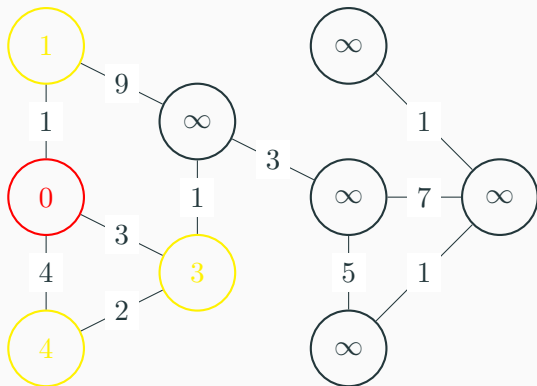


## Example

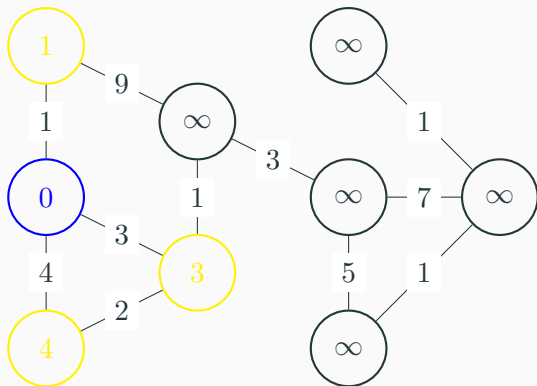




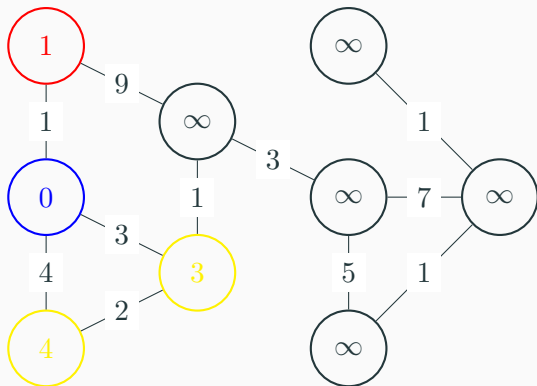
## Example




## Example



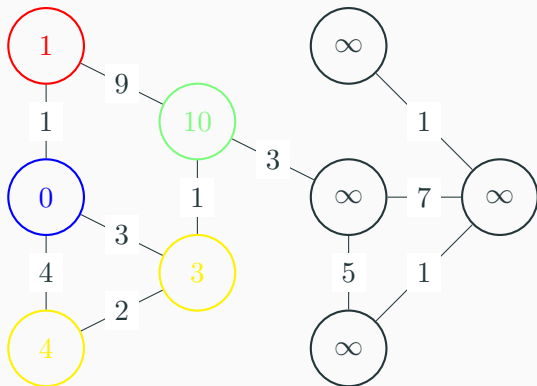
## Example



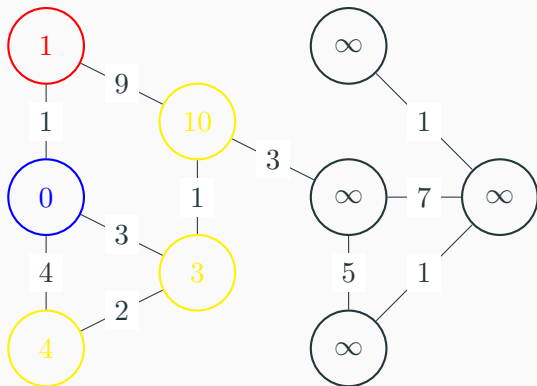
## Example



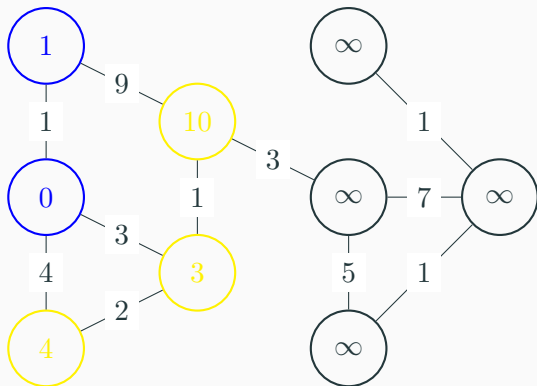
## Example



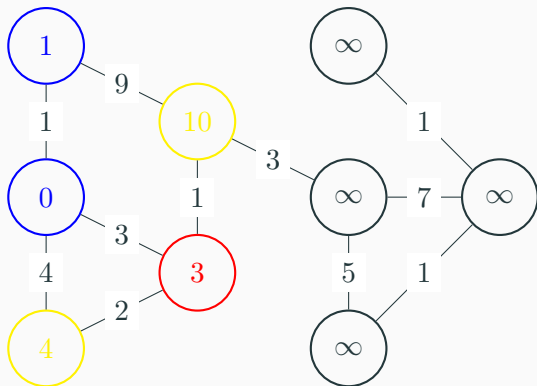
## Example



## Example

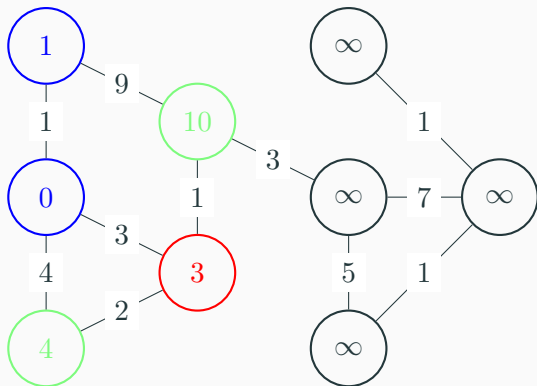


## Example

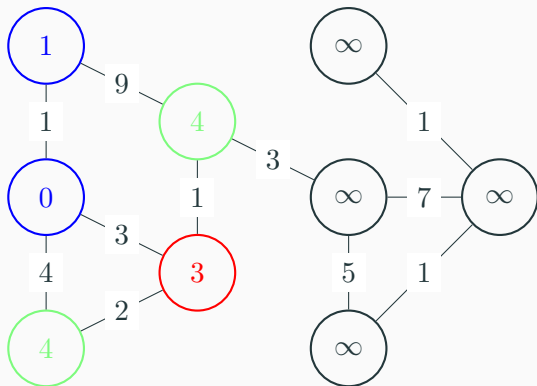




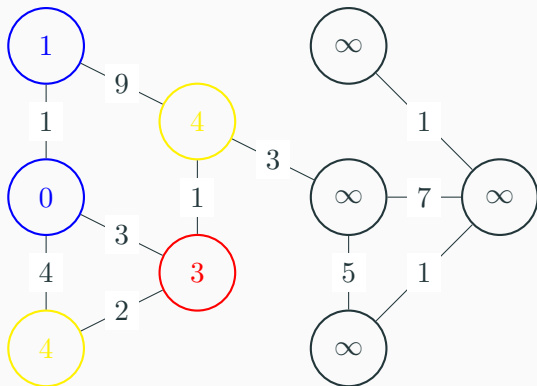
## Example



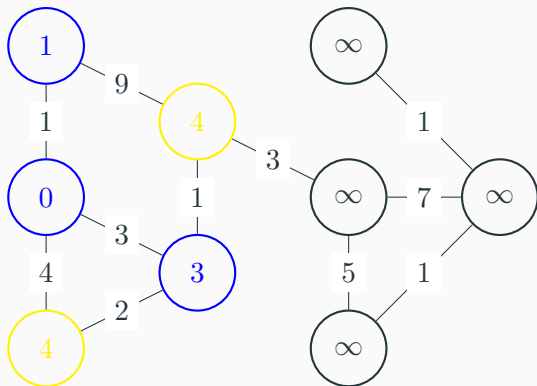
## Example



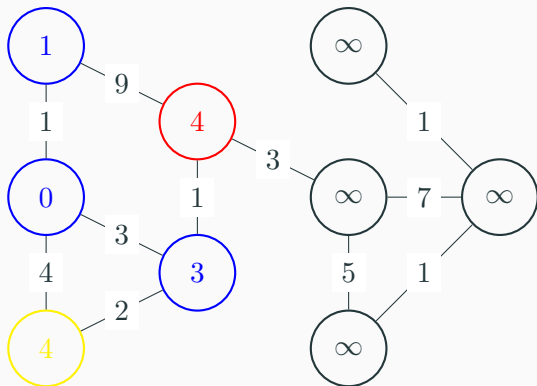
## Example



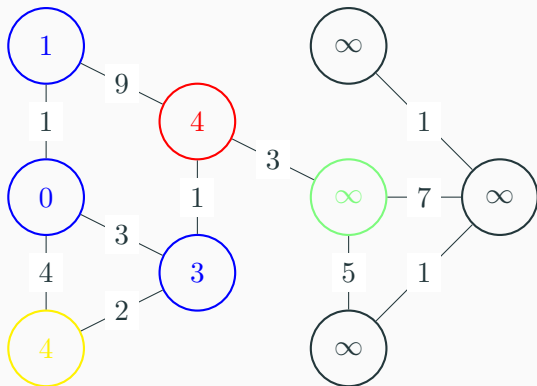
## Example



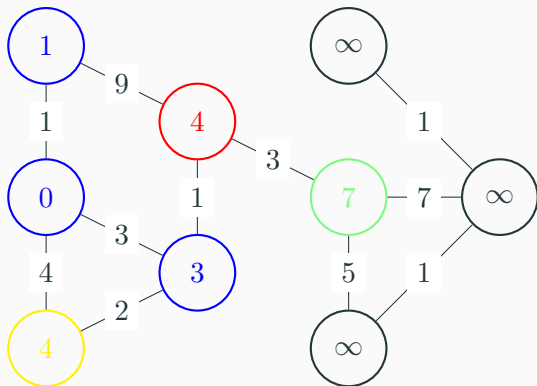
## Example



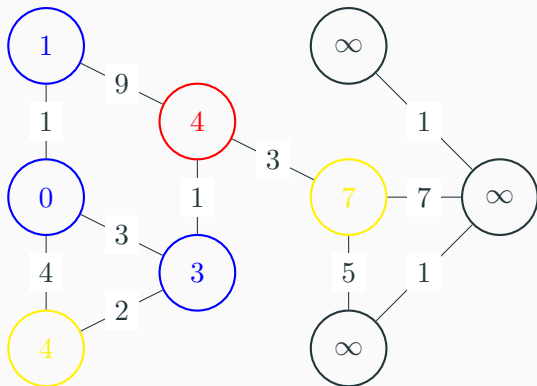
## Example



## Example

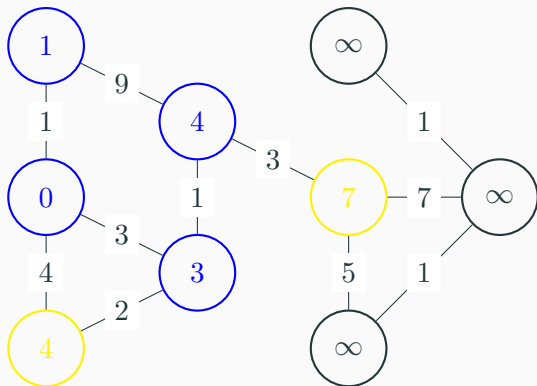


## Example

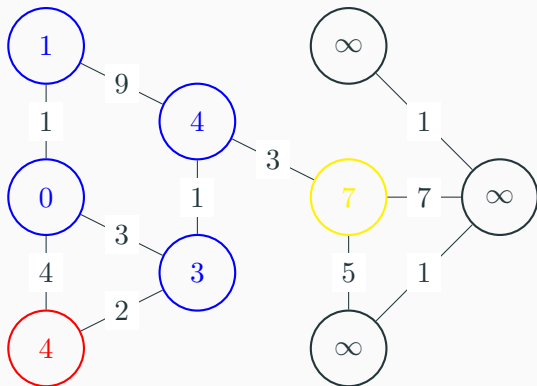




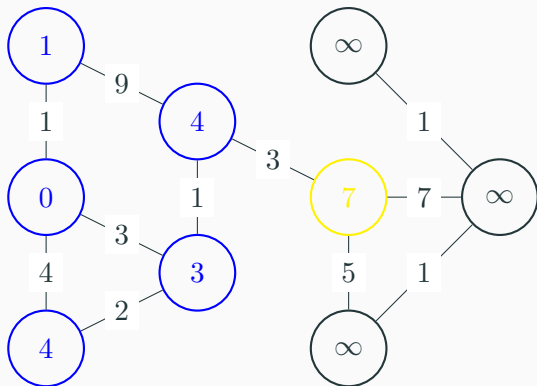
## Example



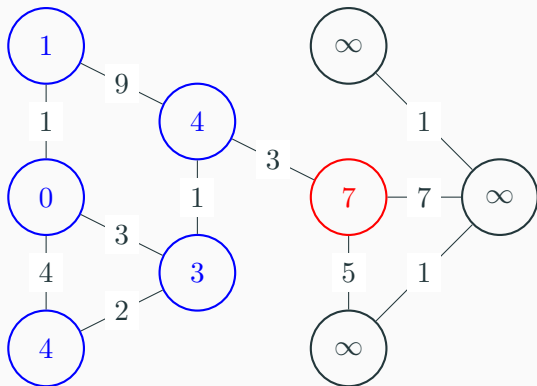
## Example



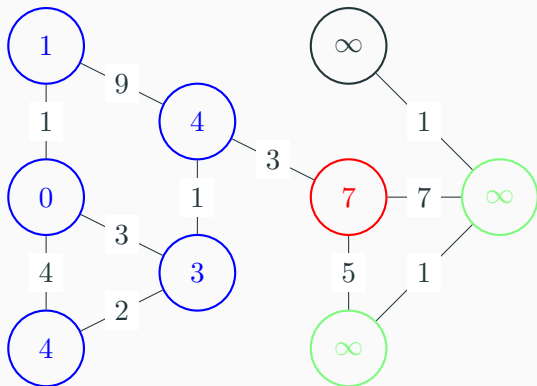
## Example



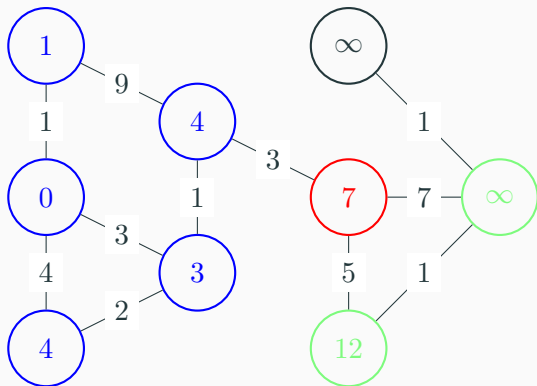
## Example



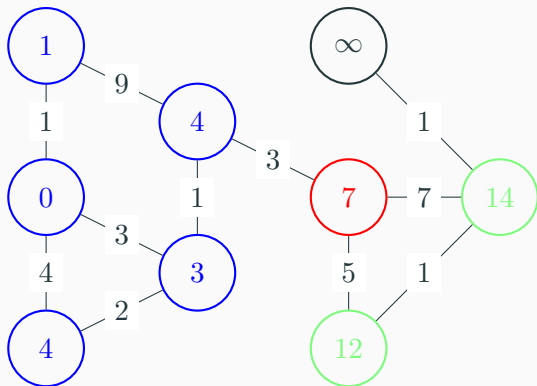
## Example



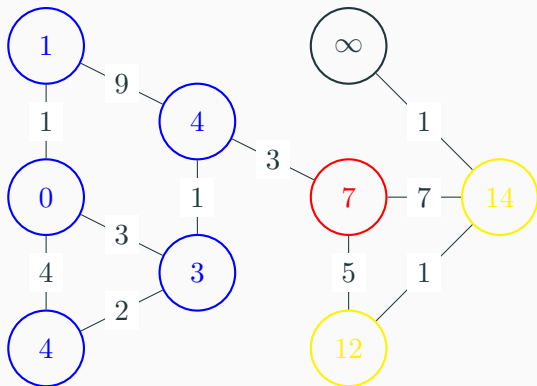
## Example



## Example

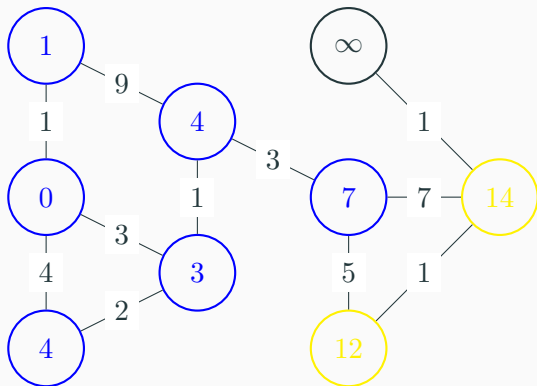


## Example

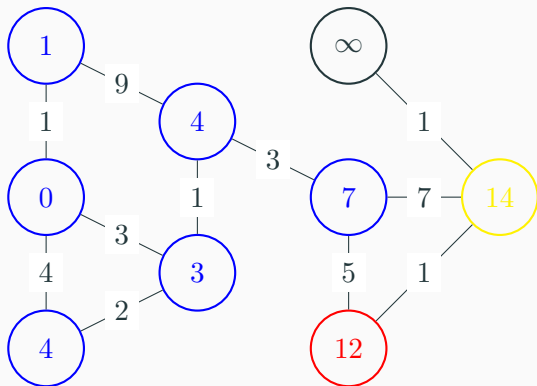




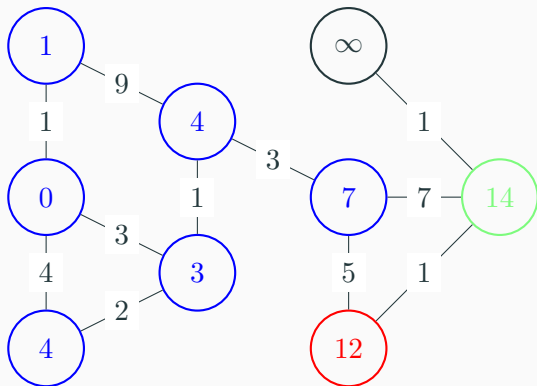
## Example



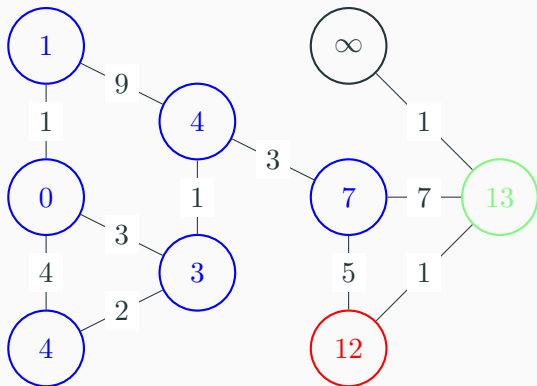
## Example



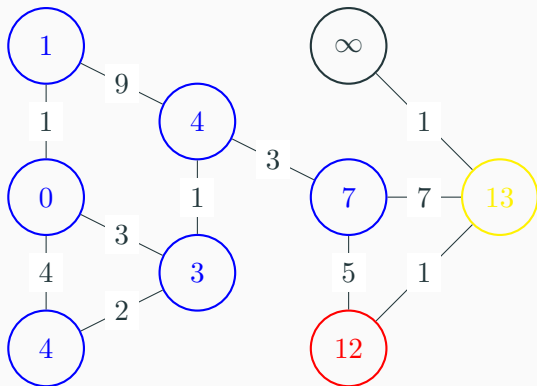
## Example



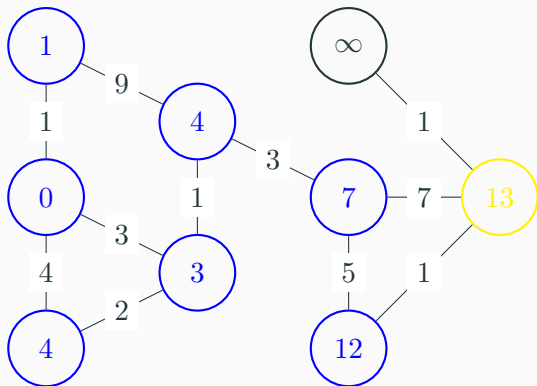
## Example



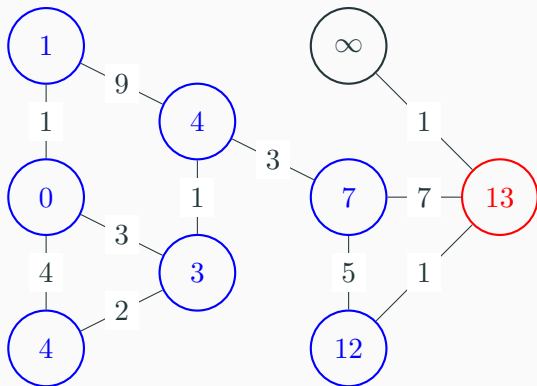
## Example



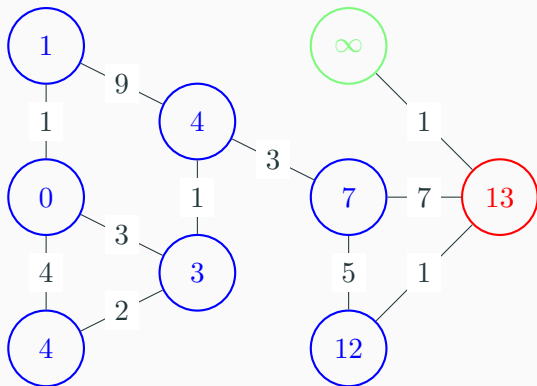
## Example



## Example

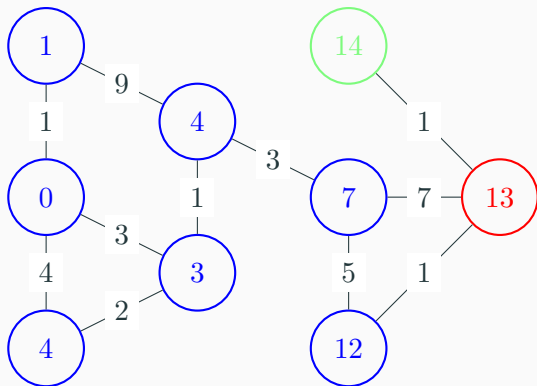


## Example

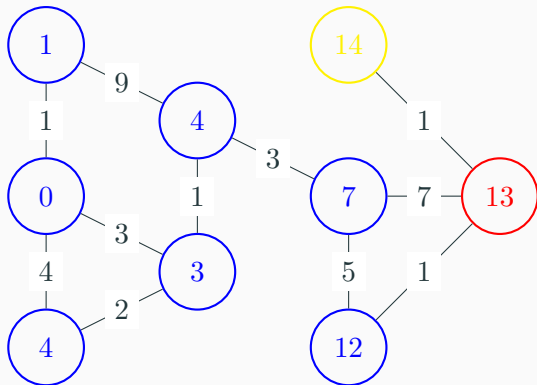




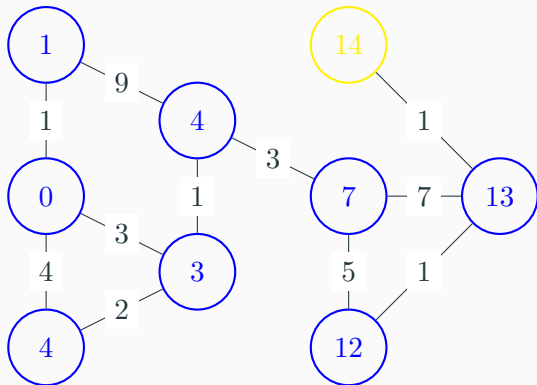
## Example



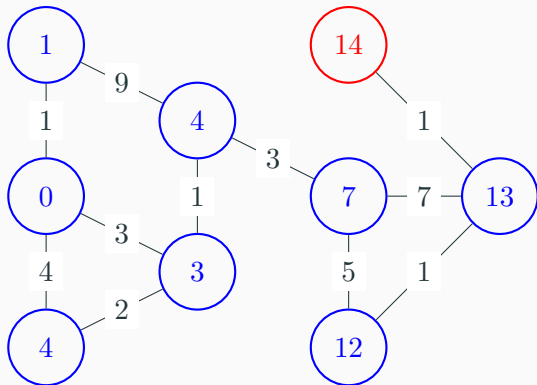
## Example




## Example



## Example



## Example



## Implementation detail

- There is one thing to watch out for
- The original version of the algorithm uses a heap where you can update the values as you change them
- This is not possible in most standard libraries, so we use a trick instead
- We push a new entry into the heap each time we update the weight to a lower one
- It is **very important** to only process an entry on the heap if the weight is correct to avoid repeated calculation when doing this!

# Dijkstra's algorithm

```
vector<edge> adj[100];
vector<int> dist(100, INF);

void dijkstra(int start) {
    dist[start] = 0;
    priority_queue<pair<int, int>, vector<pair<int, int>>, greater<pair<int, int>>> pq;
    pq.push(make_pair(dist[start], start));

    while (!pq.empty()) {
        int u = pq.top().second;
        int w = pq.top().first;
        pq.pop();
        if(dist[u] != w) continue;

        for (int i = 0; i < adj[u].size(); i++) {
            int v = adj[u][i].v;
            int w = adj[u][i].weight;

            if (w + dist[u] < dist[v]) {
                dist[v] = w + dist[u];
                pq.push(make_pair(dist[v], v));
            }
        }
    }
}
```

# Dijkstra's algorithm

- For each edge in the graph we might have to push to the heap
- We visit each vertex at most once and iterate over its neighbours
- Thus the time complexity is  $O((V + E) \log E)$
- Note that this only works for non-negative weights
- The correctness relies on the fact that once we are done with a vertex we never have to think about it again, which is not true when we have negative weights



## Dealing with negative weights

---

# Negative weights

- What do we do then when we have negative weights?
- We can use the Bellman-Ford algorithm, at the cost of a worse time complexity

# Bellman-Ford


- It is based on dynamic programming, let  $f(v, k)$  be the shortest path from the source  $u$  to a vertex  $v$  that doesn't use more than  $k$  vertices to get there
- Then as base cases we get  $f(u, 0) = 0$  and  $f(v, 0) = \infty$  for  $v \neq u$
- Otherwise we can look at  $f(w, k - 1)$  for some  $w$ , and add a single extra edge, giving the recurrence

$$f(v, k) = \min \left( f(v, k - 1), \min_{x \in N^-(v)} (w(x, v) + f(x, k - 1)) \right)$$


# Bellman-Ford

- To implement this efficiently we will make some changes from just a straight-forward DP implementation (but that would work!)
- Each  $f(v, k)$  depends only on  $f(w, k - 1)$  for some different  $w$ , so we only have to store our current row  $k$  and the last one
- This essentially gets us the usual implementation of Bellman-Ford


## Example




# Example



# Example




# Example







# Example




# Example




# Example




# Example




# Example



# Example



# Example



# Bellman-Ford

- How do we read negative cycles from this though?
- In a graph with no negative cycles a shortest path visits each vertex at most once
- Thus we can simply do  $V - 1$  more values of  $k$  to see if anything gets shorter
- If the distance to  $v$  gets shorter then there is a negative cycle on the way to  $v$



# Bellman-Ford algorithm

```
vector<edge> adj[100];  
vector<int> dist(100, INF);  
  
void bellman_ford(int n, int start) {  
  
    dist[start] = 0;  
  
    for (int i = 0; i < n - 1; i++) {  
        for (int u = 0; u < n; u++) {  
            for (int j = 0; j < adj[u].size(); j++) {  
                int v = adj[u][j].v;  
                int w = adj[u][j].weight;  
                dist[v] = min(dist[v], w + dist[u]);  
            }  
        }  
    }  
}
```

# Bellman-Ford

- This gets us a  $\mathcal{O}(VE)$  algorithm, which is quite a bit slower than Dijkstra
- A possible improvement is to store every vertex that received an update for the last  $k$  value in a queue, and only checking things they could affect at the next  $k$
- This is still worst case  $\mathcal{O}(VE)$ , but in the average case much much faster
- This is generally known as the SPFA algorithm (Shortest Path Faster Algorithm)

## All pairs shortest path

---

# Bellman-Ford

- What if we want to get the distance from every vertex to every vertex?
- If there are no negative cycles running Dijkstra's from every vertex works pretty well
- But for negative cycles running Bellman-Ford  $V$  times is pretty slow
- Luckily we can get a bulk discount in this case!
- This is known as the Floyd-Warshall algorithm

# Floyd-Warshall algorithm

- We use dynamic programming again.
- Let  $\text{sp}(k, i, j)$  be the shortest path from  $i$  to  $j$  if we're only allowed to travel through the vertices  $0, \dots, k$
- Base case:  $\text{sp}(k, i, j) = 0$  if  $i = j$
- Base case:  $\text{sp}(-1, i, j) = \text{weight}[i][j]$  if  $(i, j) \in E$
- Base case:  $\text{sp}(-1, i, j) = \infty$
- $$\text{sp}(k, i, j) = \min \begin{cases} \text{sp}(k-1, i, k) + \text{sp}(k-1, k, j) \\ \text{sp}(k-1, i, j) \end{cases}$$

# Floyd-Warshall algorithm

```
int dist[1000][1000];
int weight[1000][1000];

void floyd_warshall(int n) {
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            dist[i][j] = i == j ? 0 : weight[i][j];
        }
    }

    for (int k = 0; k < n; k++) {
        for (int i = 0; i < n; i++) {
            for (int j = 0; j < n; j++) {
                dist[i][j] = min(dist[i][j], dist[i][k] + dist[k][j]);
            }
        }
    }
}
```

# Floyd-Warshall algorithm

- Computes all-pairs shortest paths
- Time complexity is clearly  $O(n^3)$
- Very simple to code
- Though we can do better! By combining everything we did so far and mixing in an algorithm known as Johnson's algorithm, Floyd-Warshall can be beat. But this won't be in the homework (maybe the bonus problems!)